

# Resilient Distributed Datasets, Spark

*DS 5110: Big Data Systems (Spring 2023)*

Lecture 4

Yue Cheng



Some material taken/derived from:

- Matei Zaharia's NSDI'12 talk slides.

@ 2023 released for use under a [CC BY-SA](https://creativecommons.org/licenses/by-sa/4.0/) license.

## Applications

Batch

SQL

ETL

Machine  
learning

Emerging  
apps?

Scalable computing engines

Scalable storage systems



A diagram of datacenter infrastructure enclosed in a dashed rectangular border. It features two rows of server racks. The top row contains eight server racks, and the bottom row contains seven server racks. Each server rack is depicted as a black, multi-bay unit with blue horizontal light bars on its front face. The text 'Datacenter infrastructure' is centered below the top row of server racks.

Datacenter infrastructure

# What's good with MapReduce

- Scaled analytics to thousands of machines
- Eliminated fault tolerance as a concern

# Problems with MapReduce

- Scaled analytics to thousands of machines
- Eliminated fault tolerance as a concern
- **Not very expressive**
  - Iterative algorithms  
(PageRank, Logistic Regression, ...)
  - Interactive and ad-hoc queries  
(Interactive Log Debugging)
- Lots of specialized frameworks
  - Pregel, GraphLab, PowerGraph, DryadLINQ, HaLoop...



# Sharing data between stages/iterations

- Only way to share data between iterations / phases is through shared storage
  - **Slow!**
- Allow operations to feed data to one another
  - Ideally, through memory instead of disk-based storage



# Sharing data between stages/iterations

- Only way to share data between iterations / phases is through shared storage
  - **Slow!**
- Allow operations to feed data to one another
  - Ideally, through memory instead of disk-based storage
- Need the “**chain**” of operations to be exposed to make this work
- **Problem to solve:** Would this break the MR fault-tolerance scheme?
  - Retry and Map or Reduce task since idempotent

# Examples




# Examples




Slow due to replication and disk I/O,  
but necessary for fault tolerance



# Goal: In-memory data sharing



# Goal: In-memory data sharing



10-100× faster than network/disk, **but how to get FT?**

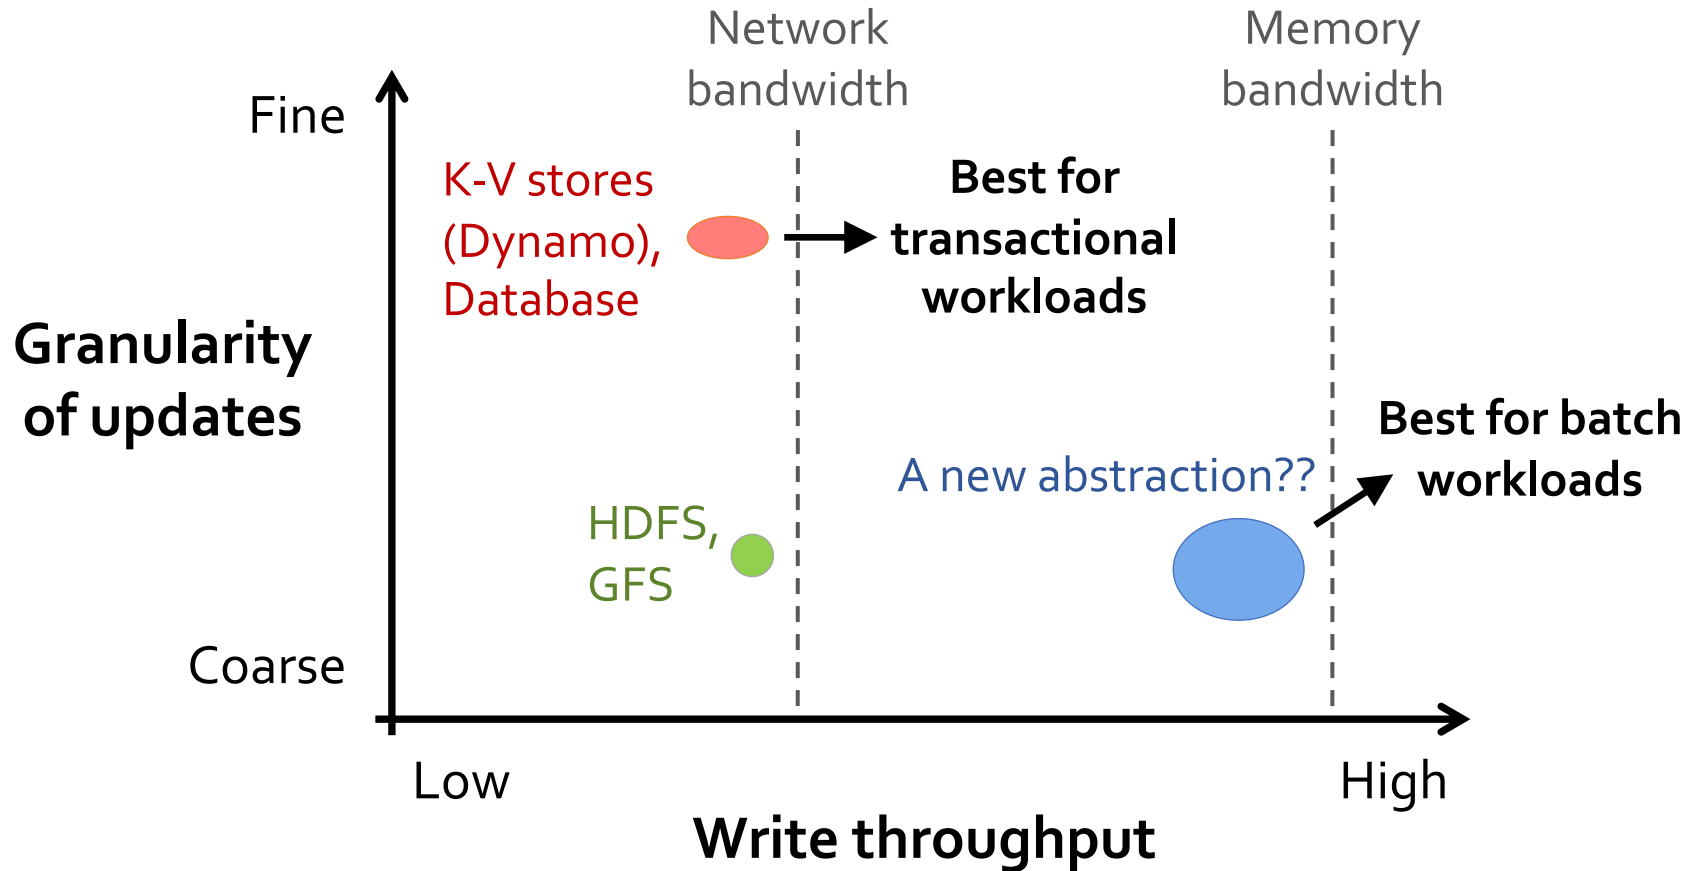
# Challenges

- How to design a distributed memory abstraction that is both **fault-tolerant** and **efficient**?

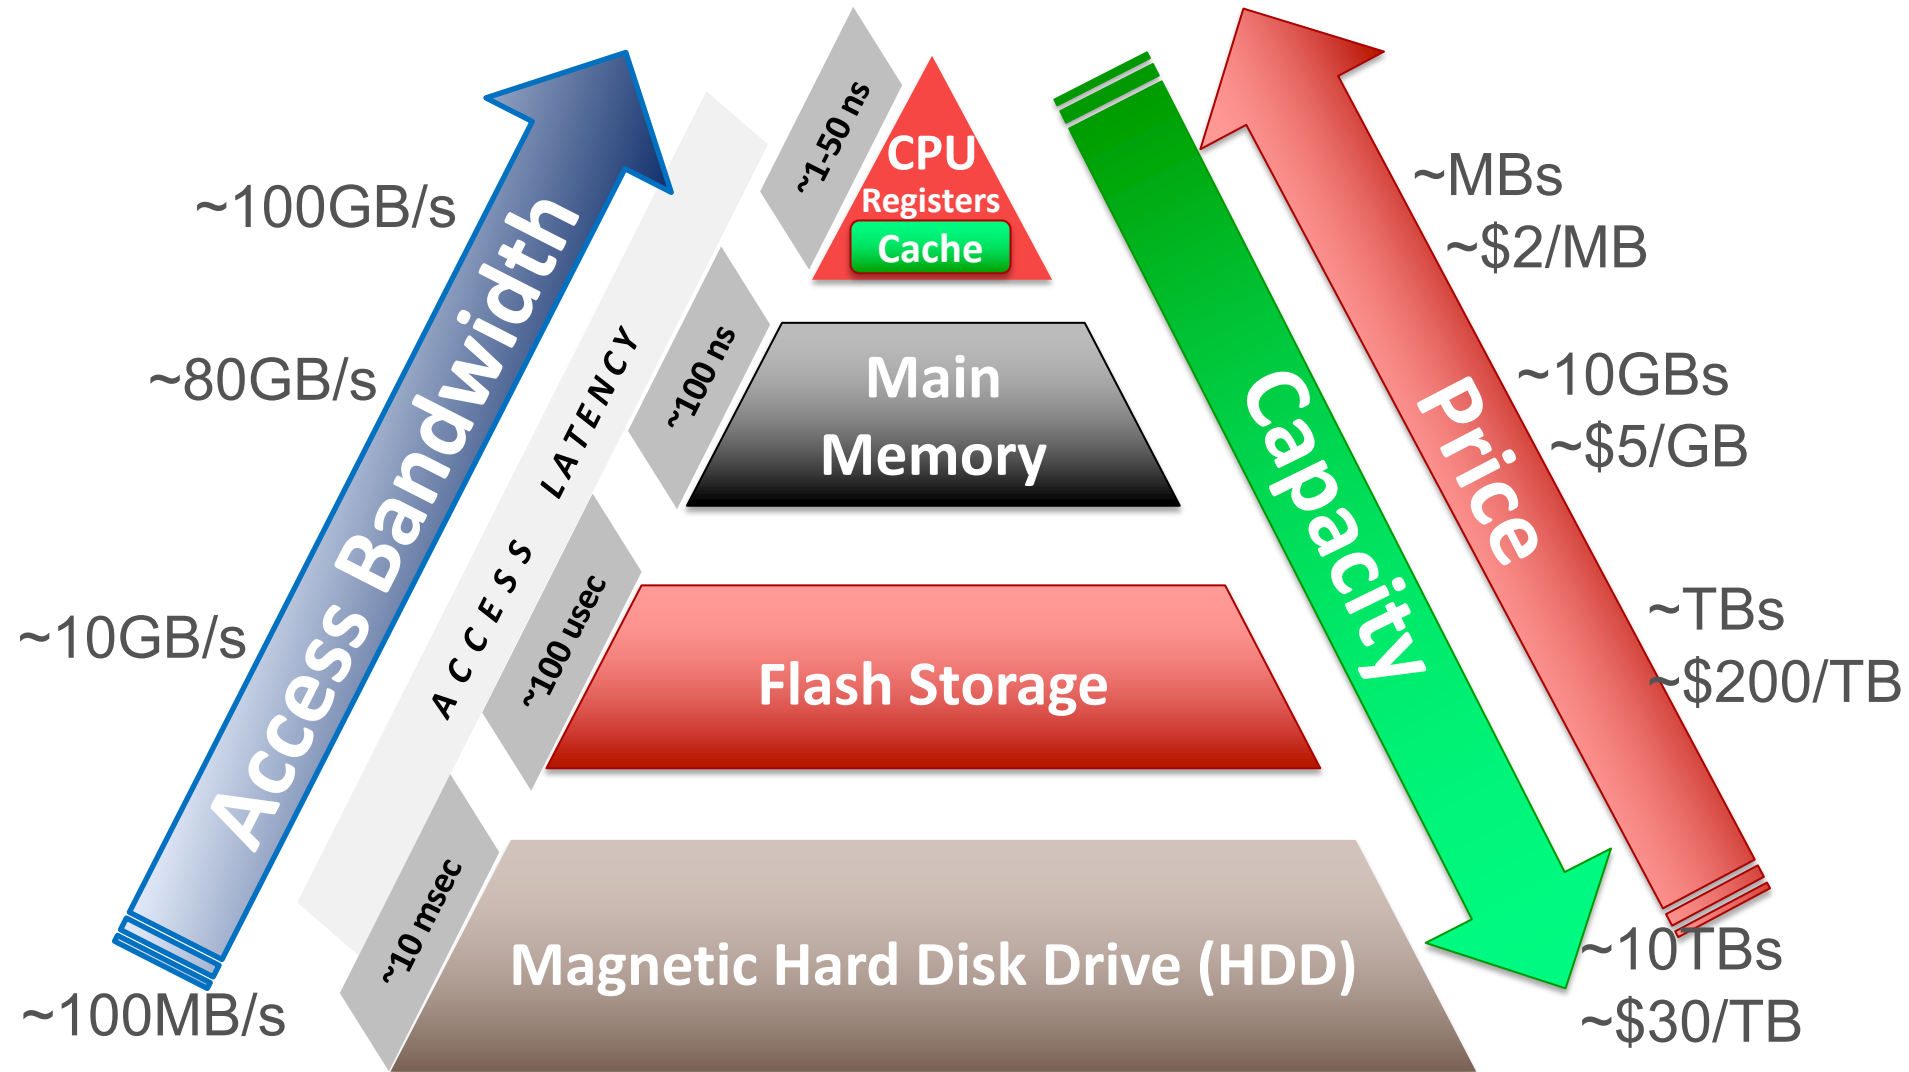
# Challenges

- How to design a distributed memory abstraction that is both **fault-tolerant** and **efficient**?
- Existing storage systems allow **fine-grained** mutation to state
  - In-memory key-value stores
  - Requires replicating data or logs across nodes for fault tolerance
    - Costly for data-intensive apps
    - 10-100x slower than memory write
  - They also require costly on-the-fly replication for mutations

# Tradeoff space



# Memory-storage hierarchy



# Challenges

- How to design a distributed memory abstraction that is both **fault-tolerant** and **efficient**?
- Existing storage systems allow **fine-grained** mutation to state

**Insight:** leverage similar coarse-grained approach that **transforms whole dataset per operation**, like MapReduce (batch processing)

- 10-100x slower than memory write
- They also require costly on-the-fly replication for mutations

# Solution: Resilient Distributed Datasets (RDDs)

- Restricted form of distributed shared memory
  - **Immutable**, partitioned collections of records
  - Can only be built through *coarse-grained*, deterministic *transformations* (map, filter, join, ...)
- Efficient fault recovery using *lineage*
  - Log **one operation** to apply to many elements
  - Recompute lost partitions on failure
  - No cost if nothing fails



# Spark programming interface

Scala API, now have multi-language bindings such as Python, Java, etc.

## Managing RDDs

- **Transformations** on RDDs ( $RDD_1 \rightarrow RDD_2$ )
- **Actions** on RDDs ( $RDD \rightarrow \text{output}$ )
- Control over RDD partitioning (how items are split over nodes)
- Control over RDD persistence (in memory, on disk, or recompute on loss)

# Transformations

Transformations (define a new RDD)	map filter sample groupByKey reduceByKey sortByKey	flatMap union join cogroup cross mapValues
---------------------------------------	---	---

RDDs in terms of Scala types → Scala semantics at workers

Transformations are **lazy operations**; cause no cluster action

# Actions

<b>Actions</b> (return a result to driver program)	collect reduce count save lookupKey
--	---

Consumes an RDD to **produce** output  
either to storage (save), or  
to interpreter/Scala (count, collect, reduce)

Causes RDD lineage chain to **get executed** on the cluster to  
produce the output  
(for any missing pieces of the computation)


# Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
errors.persist()
```

# Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
)
errors.persist()
```


```
errors.count()
```



# Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
)
errors.persist()

errors.count()
```




# Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL"))
```




# Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
```






# Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
errors.persist()

errors.count()

errors.filter(
    _.contains("MySQL")).count()
```




# Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
```




# Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
errors.filter(
    _.contains("HDFS"))
```




# Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
errors.filter(
    _.contains("HDFS"))
    _.map(_.split("\t")(3))
```




# Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
errors.filter(
    _.contains("HDFS"))
    _.map(_.split("\t")(3))
    .collect()
```




# Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR"))
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
errors.filter(
    _.contains("HDFS"))
    _.map(_.split("\t")(3))
    .collect()
```




# Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR"))
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
errors.filter(
    _.contains("HDFS"))
    _.map(_.split("\t")(3))
    .collect()
```




# Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR"))
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
errors.filter(
    _.contains("HDFS"))
    _.map(_.split("\t")(3))
    .collect()
```






# **`persist()`**

- Not an action nor a transformation
- A scheduler hint
- Tells which RDDs the Spark schedule should materialize and whether in memory or storage
- Gives the user control over reuse/recompute/recovery tradeoffs



# Lineage graph of RDDs




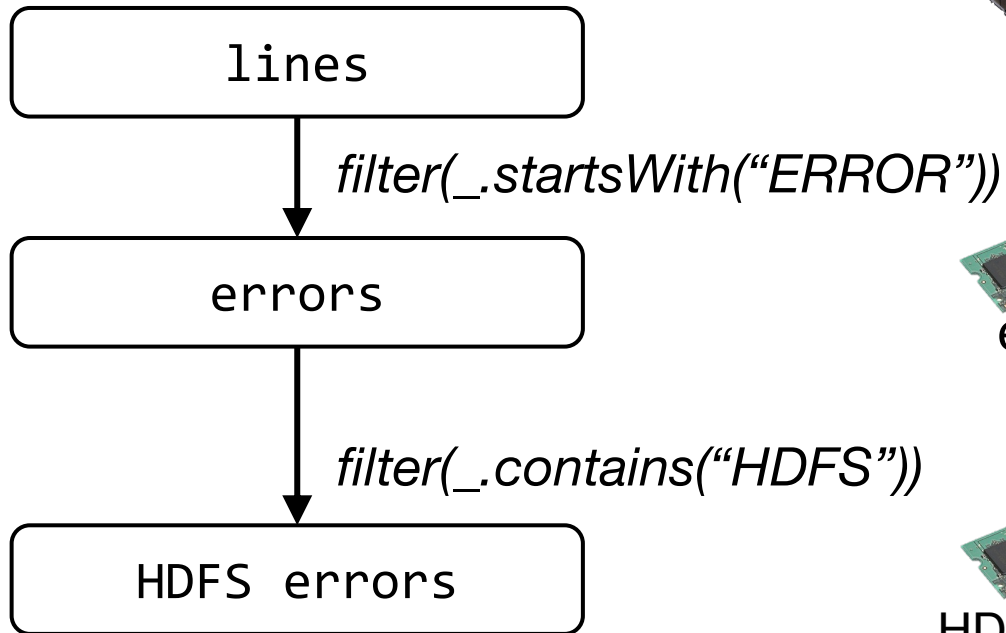
lines




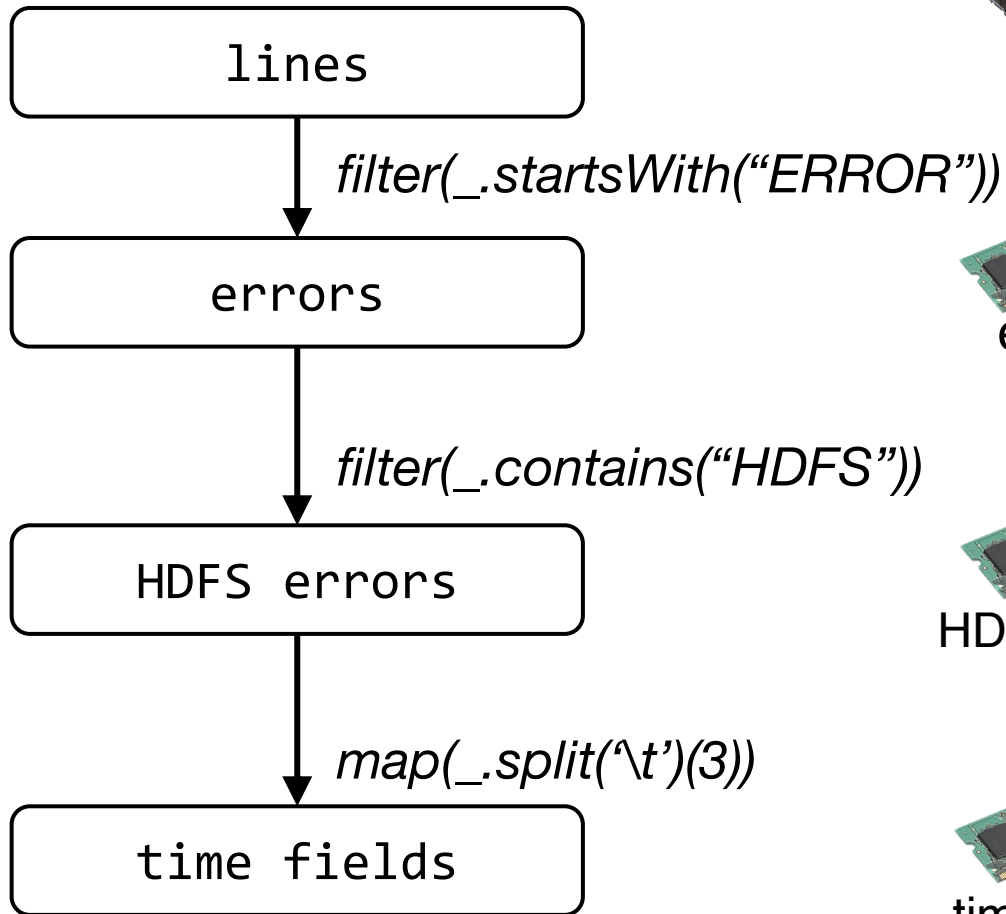
# Lineage graph of RDDs



# Lineage graph of RDDs




# Lineage graph of RDDs




# Narrow & wide dependencies

Narrow Dependencies:



Wide Dependencies:



**Narrow:** each parent partition used by at most one child partition  
(can partition on one machine)

**Wide:** multiple child partitions depend on one parent partition

Must stall for all parent data, loss of child requires whole parent RDD (not just a small # of partitions)


# Task scheduler

DAGs (directed acyclic graphs)

Pipelines functions within a stage

Locality & data reuse aware

Partitioning-aware to avoid shuffles



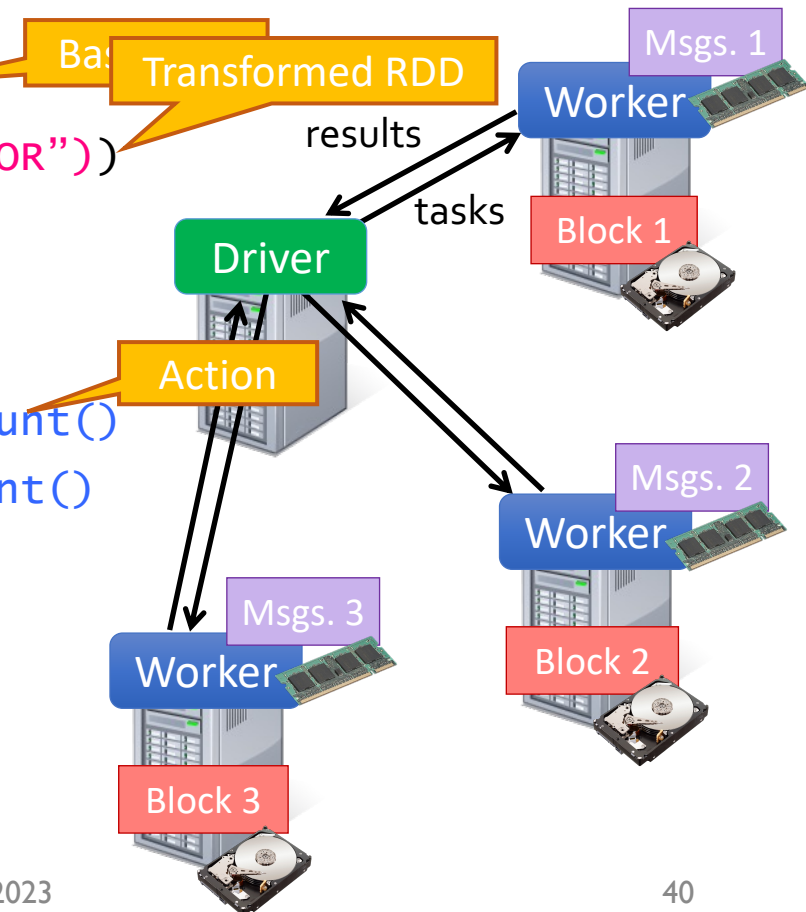
# Interactive debugging (control and data flow)

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
messages = errors.map(_.split('\t')(2))
messages.persist()

messages.filter(_.contains("MySQL")).count()
messages.filter(_.contains("HDFS")).count()
```

**Result:** scaled to 1 TB data in 5-7 sec  
(vs 170 sec for on-disk data)






# Fault recovery

- RDDs track the graph of transformations that built them (their *lineage*) to rebuild lost data

E.g.:

```
messages = textFile(...).filter(_.contains("error"))  
                        .map(_.split('\t')(2))
```




# Fault recovery

- RDDs track the graph of transformations that built them (their *lineage*) to rebuild lost data

E.g.:

```
messages = textFile(...).filter(_.contains("error"))  
                        .map(_.split('\t')(2))
```




# Fault recovery


- RDDs track the graph of transformations that built them (their *lineage*) to rebuild lost data

E.g.:

```
messages = textFile(...).filter(_.contains("error"))  
                        .map(_.split('\t')(2))
```



# Fault recovery results



# Example: PageRank

1. Start each page with a rank of 1
2. On each iteration, update each page's rank to

$$\sum_{i \in \text{neighbors}} \text{rank}_i / |\text{neighbors}_i|$$

links = // RDD of (url, neighbors) pairs

ranks = // RDD of (url, rank) pairs

```
for (i <- 1 to ITERATIONS) {  
  ranks = links.join(ranks).flatMap {  
    (url, (links, rank)) =>  
      links.map(dest => (dest, rank/links.size))  
  }.reduceByKey(_ + _)  
}
```

# Example: PageRank

1. Start each page with a rank of 1
2. On each iteration, update each page's rank to

$$\sum_{i \in \text{neighbors}} \text{rank}_i / |\text{neighbors}_i|$$

RDD[(URL, Seq[URL])]

links = // RDD of (url, neighbors) pairs

ranks = // RDD of (url, rank) pairs ← RDD[(URL, Rank)]

```
for (i <- 1 to ITERATIONS) {  
  ranks = links.join(ranks).flatMap {  
    (url, (links, rank)) =>  
      links.map(dest => (dest, rank/links.size))  
  }.reduceByKey(_ + _)  
}
```

For each neighbor in links emits (URL, RankContrib)

Reduce to RDD[(URL, Rank)]


# Join (⋈)

Alice	5	⋈	Alice	F	=	Alice	5	F
Bob	6		Bob	M		Bob	6	M
Claire	4		Claire	F		Claire	4	F

A	5	⋈	C	5
A	2		B	2
A	3		A	3
B	4		B	4
B	1		A	1
C	6		B	6
C	8		C	8

If partitioning doesn't match, then need to reshuffle to match pairs. Same problem in `reduce()` for MapReduce.


# Optimizing placement



- Links & ranks repeatedly joined
- Can *co-partition* them (e.g. hash both on URL) to avoid shuffles
- Can also use app knowledge, e.g., hash on DNS name
- `links = links.partitionBy(new URLPartitioner())`



# Optimizing placement



- Links & ranks repeatedly joined
- Can *co-partition* them (e.g. hash both on URL) to avoid shuffles
- Can also use app knowledge, e.g., hash on DNS name
- `links = links.partitionBy(new URLPartitioner())`

Q: Where might we have placed **persist()**?


# Co-partitioning example

Co-partitioning can avoid shuffle on join

But, fundamentally a shuffle on **reduceByKey**

Optimization: custom partitioner on domain

# PageRank performance



\* Figure 10a: 30 machines on 54 GB of Wikipedia data computing PageRank