

# Reinforcement Learning Systems: Ray

*DS 5110: Big Data Systems (Spring 2023)*

Lecture 8a

Yue Cheng



## Applications

Batch

SQL

ETL

Machine  
learning

Emerging  
apps?


Scalable computing engines

Scalable storage systems


A diagram showing a datacenter infrastructure. It consists of two rows of server racks, each row containing eight racks. The racks are black with blue horizontal lines and small blue lights at the bottom. The entire set of racks is enclosed in a dashed rectangular border.

Datacenter infrastructure


# Reinforcement learning



Atari breakout



Pong: after 30  
mins of training





Pong: DQN  
wins like a boss

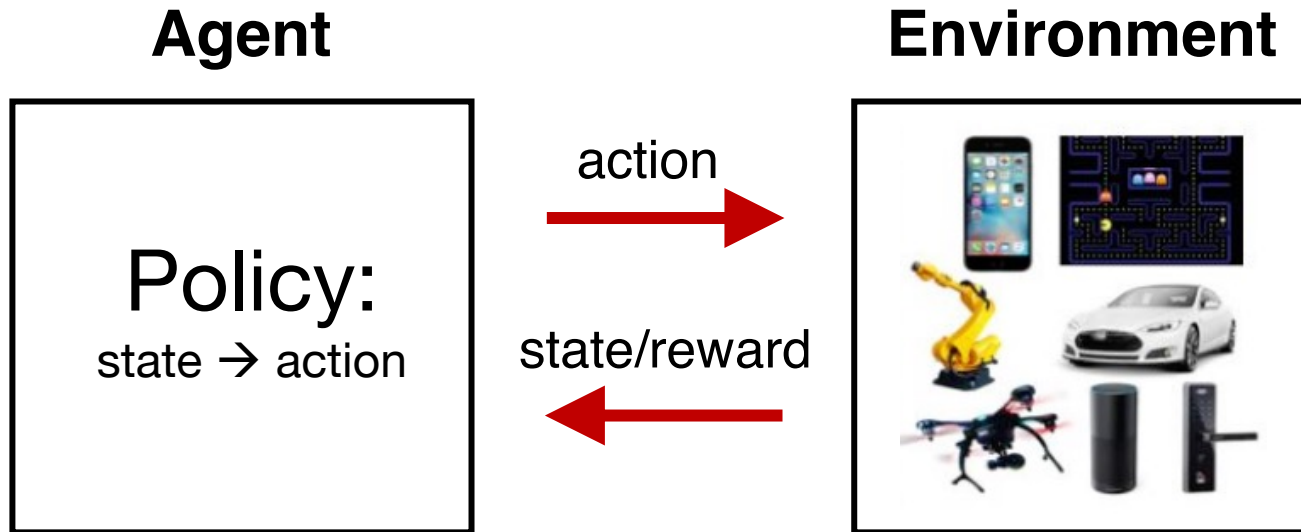
\*: Playing Atari with Deep Reinforcement Learning: <https://arxiv.org/abs/1312.5602>

# RL application pattern

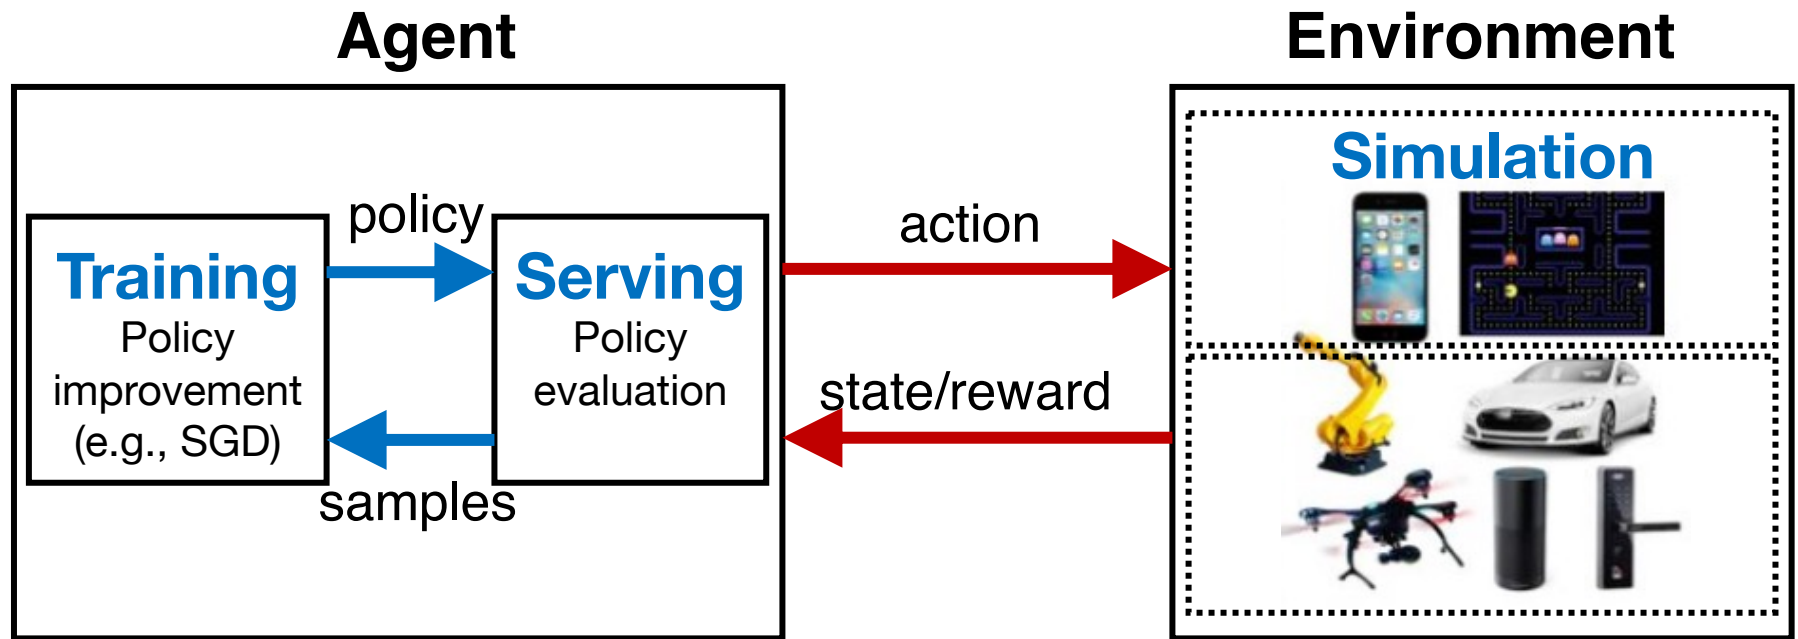
- Process inputs from **different** sensors in **parallel & real-time**
- Execute large number of simulations, e.g., up to 100s of millions



# RL setup



# RL setup in more detail



# RL application pattern

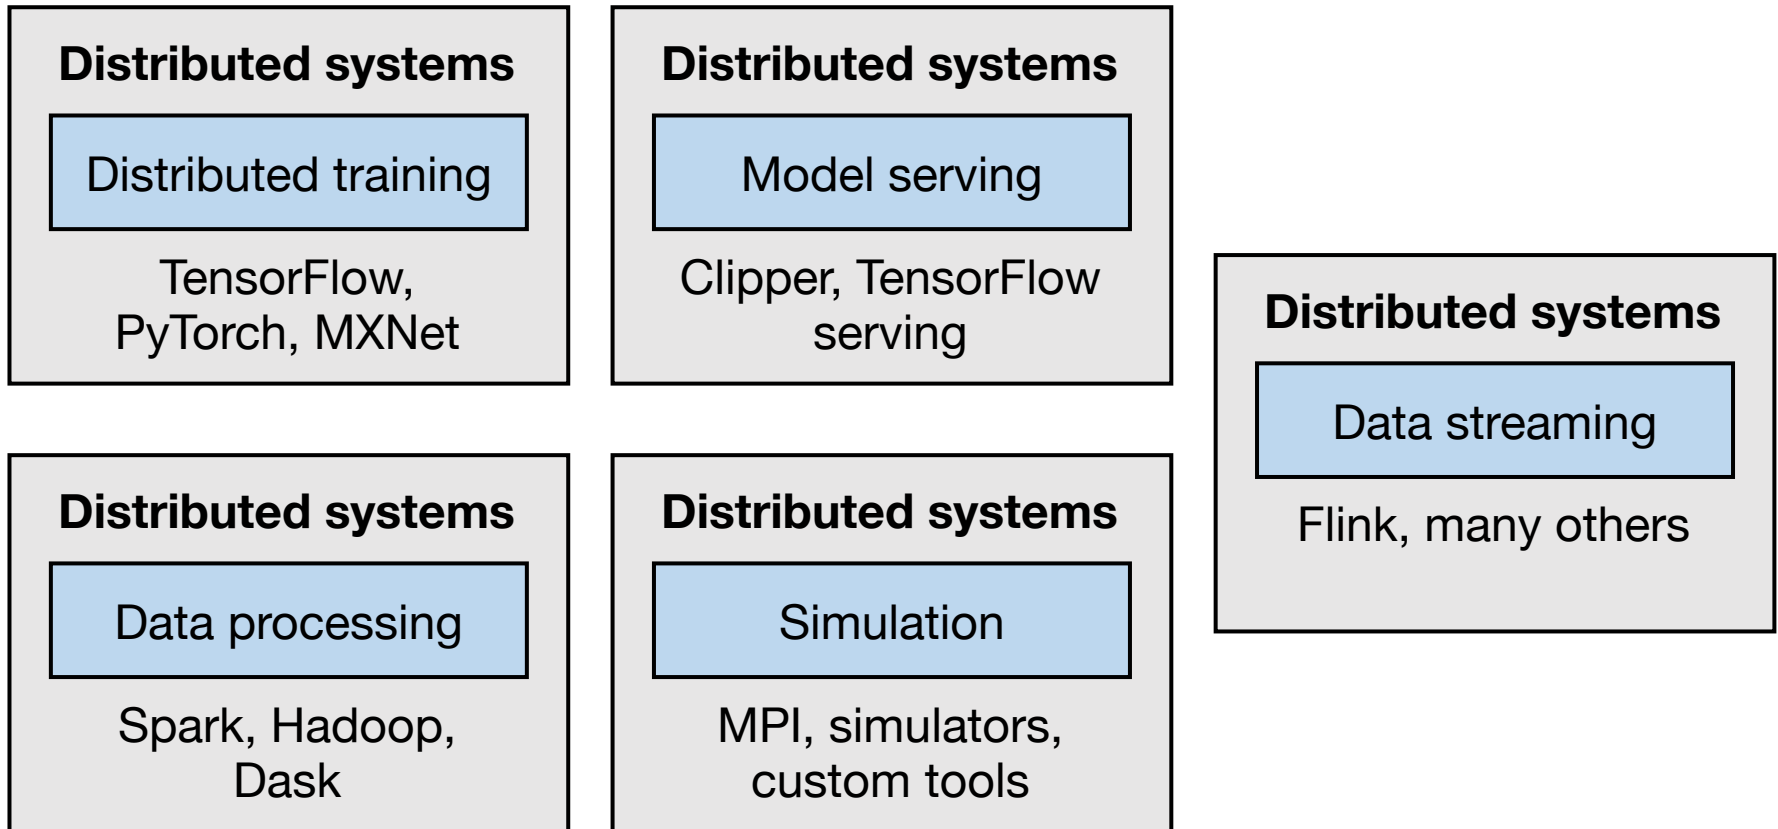
- Process inputs from different sensors in parallel & real-time
- Execute large number of simulations, e.g., up to 100s of millions
- Rollouts outcomes are used to update policy (e.g., SGD)

# RL application requirements

- Need to handle dynamic task graphs, where tasks have
  - Heterogeneous durations
  - Heterogeneous computations
- Schedule millions of tasks / sec
- Make it easy to parallelize ML algorithms (often written in Python)



# The ML/AI/data ecosystems today



Emerging AI applications require **stitching**  
together **multiple** disparate systems

Ad hoc integrations are **difficult to manage and program!**

# Ray API

## Tasks

```
futures = f.remote(args)
```

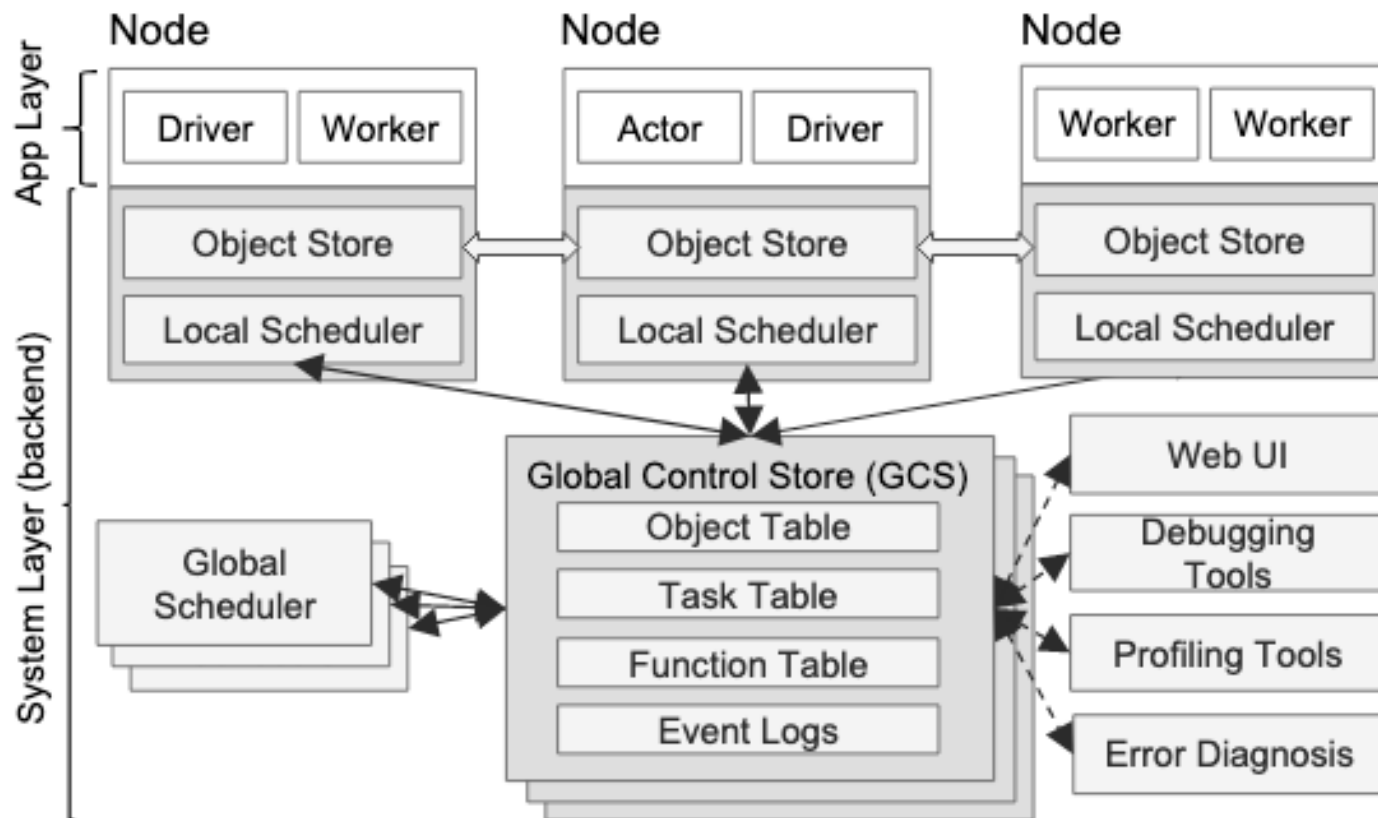
## Actors

```
actor = Class.remote(args)  
futures = actor.method.remote(args)
```

```
objects = ray.get(futures)  
ready_futures = ray.wait(futures, k, timeout)
```

# Ray API examples: Demo

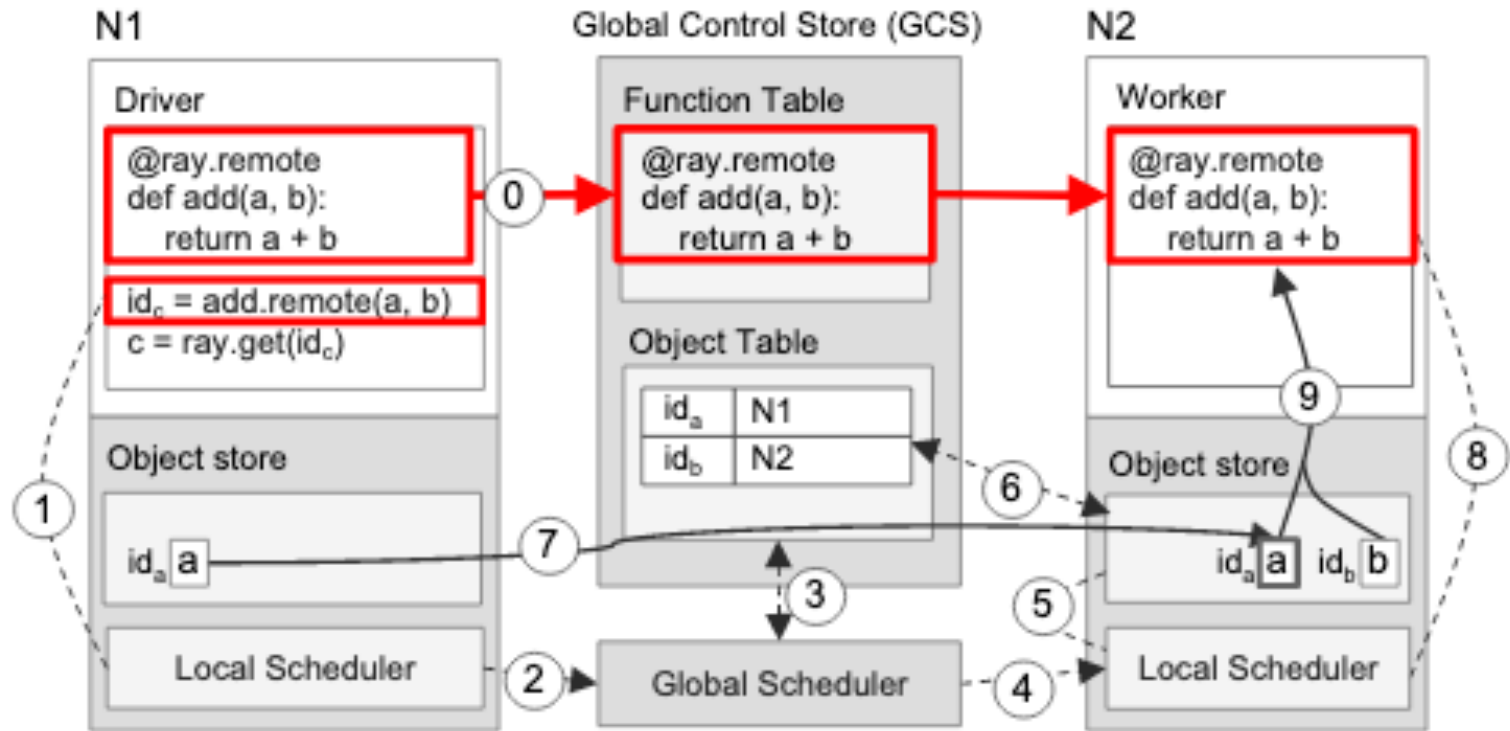
# Ray architecture



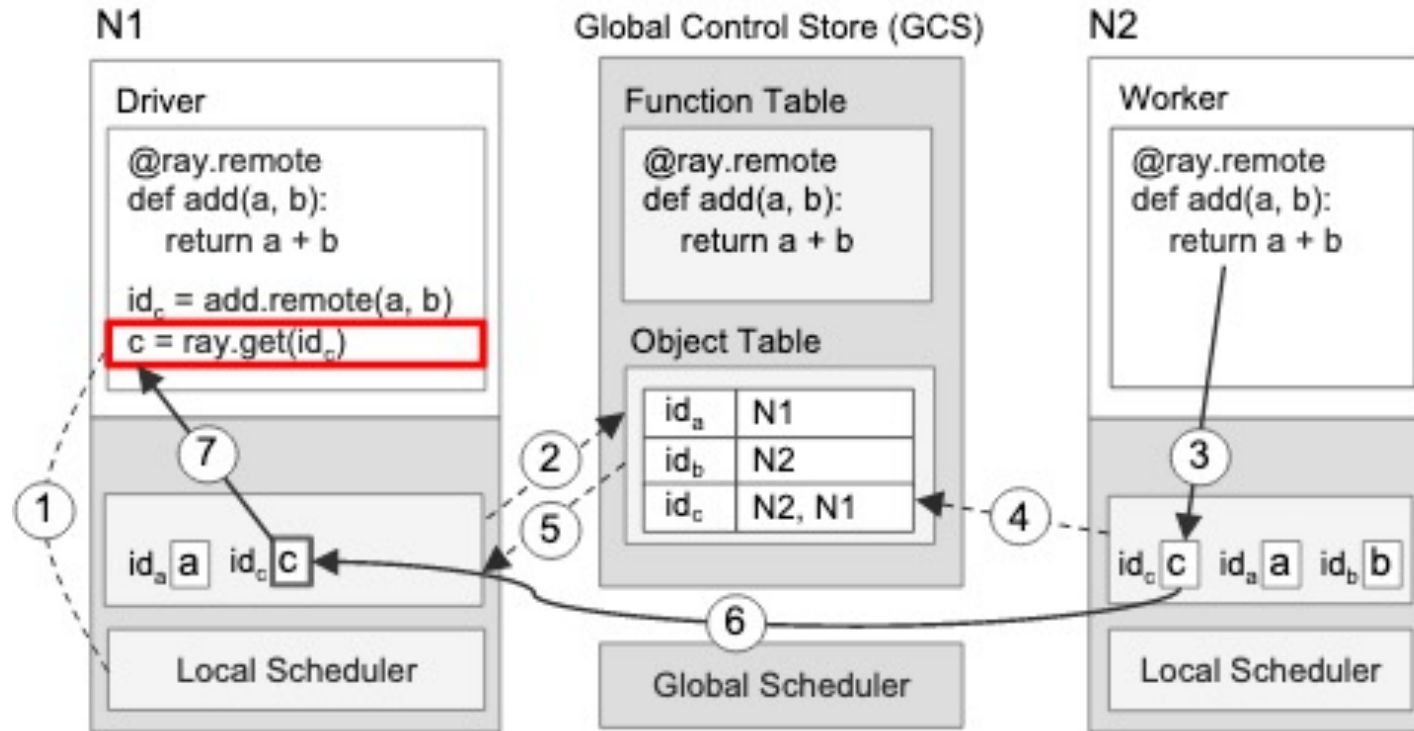
# Global control store (GCS)

- Object table
- Task table
- Function table

# Executing a task remotely



# Returning the results of a remote task





# This Wed: Federated Learning Systems